

ChatScript Client/Server Manual

Copyright Bruce Wilcox, gowilcox@gmail.com brilligunderstanding.com Revision 7/18/2021 cs11.5

- Running the server
- Unique User Names
- ChatScript protocol
- Communicating with the Server
- Testing the server
- Revising a live server
- Revising a topic
- Preparing for compiling on a server
- Testing for server presence
- Server crashes and cron
- CPU vs IO bound
- Memory issues with multiple servers on a machine
- Commands affecting the server
- Command Authorization
- RESTful server
- Encryption

While the system defaults to running as a stand-alone chatbot under Windows, when run under LINUX it defaults to being a server.

Nominally (meaning depending on hardware and what your bot does) ChatScript can process a volley on a single core in 10 milliseconds on a slow machine, thus handling 100 volleys every second from different users using one core. A human-human volley is often around 15 seconds, so handling 1000 simultaneous users with a single core slow server is not unreasonable.

The fastest server OS for ChatScript is Linux. The Mac tends to misfire in the OS itself with heavy client loads. Windows is a much slower server in general. And Linux version of CS has support for forking chatscript (no such support under Windows), so you can run a fork of the engine on every core, saturating cpu processing to the max while still serving a single port. Speedup is nearly linear per core added.

Running the Server

When you run the Mac/LINUX program, it defaults to server mode, port 1024. To run the server under Windows you must give it a command line parameter specifying a port. There are various command line parameters to affect behavior, described at the end.

Unique User Names

ChatScript maintains an independent history with each user-bot combination in a single file in the USER's directory. It is nominally up to you to define some unique name for each user. There is no login validation service provided by ChatScript; that is your responsibility. Some simple things which ChatScript supports directly are:

1. If the login name is “.” , the system will assign a user name of the IP address it receives. This doesn't work from localhost nor will it work if you have a server of your own relaying between the client and the ChatScript server (since the IP address will always be of your relay server).
2. If the login name is *guest*, the system will assign a user name of guest concatenated with the IP address. You could also use the user's email address as a login id. ChatScript will automatically convert periods and @ in a user's name to _ , so a login like *gowilcox@gmail.com* will become *gowilcox_gmail_com*. Likewise logging in as guest will result in something like *guest_123_124_155_12*. Of course, if the user eventually comes back some other day on a different IP address, they lose their history. Can't be helped. And you end up with dead files.

Communicating with the Server

The client webpage/program connects on a socket to the IP and port of the server. If you are coming from a webpage, the webpage must establish the socket. The system does not use HTTP. HTTP is an agreement on what port to use (a standard http port) and what message protocols look like to the HTTP server. Similarly ChatScript uses an agreement on what port to use (but you get to specify the port) and what the message protocols look like.

And, each communication is a one-shot deal. The socket is made, the client sends a message to the server, the server sends data back, AND CLOSES THE CONNECTION. Given that potentially thousands of users may be using the server at the same time, it is undesirable for it to try to maintain that many open sockets. If for some reason you need to maintain a permanent connection to a client, you can write an intermediary server program that has a permanent socket to the client, and relays messages back and forth between the client and the ChatScript server.

This is also what you would do if you wanted ChatScript to “push” unsolicited messages to the client. Your intermediary server can use timeouts to decide to send a specific message to the ChatScript server, and then relay the result back to the client as an unsolicited output. Similarly, your intermediate server might receive various asynchronous events and signals, and can pass requests to ChatScript at that time and pass the result back to a permanently connected client.

Each chat volley (incoming and outgoing message) is an entirely different connection. This means the Chatbot server is not devoting resources waiting for a user to reply and doesn't care how long it takes the user to come back again.

ChatScript Protocol

The message a client sends is a concatenation of three null-terminated text strings. Or in place of the first two null-termination markers, you may use Ascii 1 (since JavaScript doesn't play well with trying to join null-terminated strings) so that you send a single null-terminated string with ascii 1 as separators between components. The first component is the user login name. The second component is the name of the chatbot to talk to. If this is a null string (or just Ascii 1), the system will connect to the default bot.

The third component is the message. If the message is null, this is a start of new conversation. This **MUST** be the first thing you do with a new user. Ideally you do it whenever a new conversation is starting with that user which is how the system knows the old conversation ended. Usually script will detect that this is the start of a new conversation and say something like *hello* or *welcome back* to indicate the two parties are starting up a new conversation, though through the history file the system may have a lot of information about what has gone on in prior conversations. As long as the user is connected to the webpage, for example, you wouldn't send a startup message again.

Due to the requirement of a unique user name, you **NORMALLY** require the user to enter a login name once on the client, after which you pass that on each transmission to the server. You can bypass asking for a user name if you always just use the "guest" or "." user names.

The message sent to the server during a conversation should never be null (since that looks like a conversation start). Either always prepend a blank on every line from the user, or add a blank if the user presses ENTER without anything else or pass along the newline/cr character.

The chatbot can wait forever for each input (the connection is terminated for each volley) and the only way to know that the human "left" is when the human "comes back" with a start of a new conversation.

Testing the Server

You can test everything on your own machine in Windows using batch files in **SERVER BATCH FILES**. Launch the server by double clicking on server.bat and then launch the client by double clicking on localclient.bat. Since this is on your own machine, firewall opening a port is unnecessary. Remember that to be a remote server you need to make your port available for inbound TCP if you have a firewall - if you don't have a firewall you must be insane).

If you run the Windows ChatScript engine with `client=ip:port` as an argument, it will act as a client to talk to and test a remote server. The ip address must be numeric and the `:port` is optional.

The client will start a new conversation and then loop with you conversing to it (assuming your server is running). If you issue a `:restart` command to the server, the client will prompt you for a new login. If you prefix your login name with `*`, then you get to speak first to the chatbot continuing any prior conversation (normally the chatbot speaks first starting a new conversation).

For LINUX, just perform the equivalent commands of the batch files (except that since you can't readily run multiple apps, you'll have to background the server presumably by doing a `nohup` command on it.

Revising a Live Server

When you want to change the contents of the server, obviously you could just stop and start it again. Of course that blocks people from using it in the interim and you might even cut off a user from getting his response. You don't have to do that.

You can restart a live server without interfering with people. In actuality, the server starts up by loading all of its data into memory (except for specific users). So the dictionary, topic data, livedata, etc are all memory resident. This means you can, while the server is running, revise those files on disk. You can, for example run a stand-alone copy and rebuilt topics from raw data. Or you can edit livedata files or dictionary files. None of this impacts a running server.

Then, as an authorized user, you can issue a `:restart` command. Because the chatbot server handles one user one volley at a time, when you are talking to the server no one else is (briefly). So you can tell the server to reload all of its data and that will complete your turn and it will then merrily handle the next user with the new data.

Revising a topic

Each user's record tracks where they have been in a topic. If you have changed that topic, the record becomes invalid and any memory it has of where they have been in that topic, or if they are sitting at a rejoinder there waiting for the next user input, will be forgotten. It's now a fresh topic. All other unchanged topics will not affect the user in any way.

The user's topic data does not have the actual rules of the topic (it is not a copy of the topic) because that would make the files huge. Instead, the checksum of the topic is stored with the user's topic data in `USERS` when they start using up rules from a topic. And each top level rule is identified by a corresponding

bit in the user's topic data, which can be compressed if a block of rules is all used or all unused. So effectively we turn on a bit whenever a rule is used up.

When you bring a new topic into the system, it has a different checksum. Therefore the system detects that the user's file does not match the new topic data and should be considered invalid. IF the checksum of the new topic is 0 (a special value which you can request with the flag "safe" on the topic) then the system considered the topic automatically compatible. So if you just change spellings or add rules at end or add rejoinders (which are not top level), such a topic can continue to be safely used with a user's saved data.

In reality, the only change where you have to worry about compatibility is altering gambits, since you can always safely add responders at the end or rejoinders anywhere. And you can even add gambits using a trick. You could insert a block of n gambits to replace an existing gambit via making the block be a new topic and replacing the existing gambit with a call to that new topic. This is a "safe change" in that someone who has not reached that gambit will get the new experience. Of course some one who has already been to that gambit will not see the revision, but can continue on the rest of the topic as it used to be.

And technically I speak a mistruth in saying that you can just dump responders at the end. IF you use a segmented responder to access the topic `^respond(~yourtopic.secondhalf)` then you would be impacting that. You can get around that by instead replacing a responder with `u: () ^respond(~auxresponses)`.

Preparing for compiling on the Server

I develop the source on a Windows machine and transfer it to a LINUX box. To insure the source does not have carriage returns, I use `:clean` to read and write all src directory files without carriage returns.

Testing for server presence

To truly know if the server is alive, you actually need to send it a message that it responds to. This means implementing **something** in your bot that reacts rapidly to your message. The fastest messages are OOB ones, because they do not parse. So early in your control script you can do something like:

```
u: ( \[ ping \]) Pong.
```

and use that to prove the server is fully functional.

Server Crashes & cron

If the server crashes, it may automatically recover, generating the message

Sorry. I forgot what I was thinking about.

If your bot personality assigns a message to the user variable \$crashmsg, it will use that instead. One cannot guarantee the server doesn't go down completely, and I recommend it be on a cron job trying to start it maybe every minute or every 5. The system will detect if its port is already busy and not start a new copy if the old one is still running.

Be advised that ChatScript assumes the current directory is the one the executable is in and accesses its data relative to that. For cron this means you want an entry like:

```
0,5,10,15,20,25,30,35,40,45,50,55 * * * * cd /home/bruce/ChatScript;  
./ChatScript/LinuxChatscript32 2>/home/bruce/cronserver.log
```

Logging

CS user or server logs are never written to anything other than local filesystem. A server log could in theory be written remote to a database since it is done as a single log line, but user logs can involve tracing, which must go to local filesystem since the log is written incrementally within a line and there would be inordinate network traffic.

Still, for server logs it is a constantly increasing record count and performance would suffer significantly so logs go to the local filesystem and you have to get logs from the local machine(s). This is what they do at Amazon and they have tools to yank logs from all deployed servers.

CPU vs IO bound

The ChatScript server consists of a main thread to handle computing a response, a thread to accept incoming connections, and threads for each connection spawned. The connection threads handle reading the message from the user, getting the attention of the main thread to get a response, and then passing that response back to the user. Any logging is normally done from the connection threads, so the main server is free to spend all its effort handling volleys.

Under a full load of users, the main thread will nominally be always busy and the system CPU bound as a consequence. However, the main thread must also read and write user data for each volley. That may slow things down and make the system IO bound if there are lots and lots of user files around and the system is using cloud-based files instead of local ones. This will lower throughput significantly. It can be compensated for by locally caching active user data. There is a parameter to tell the server to track some number of users in memory. It will write out those memory copies periodically, but obviously if the server crashes, some users may be 50 volleys out of date.

Memory Issues with multiple servers on a machine

If you run multiple servers on a machine, you may find one of them fails to start because it runs out of memory. ChatScript allocates all the memory it thinks it needs at startup, so if it succeeds, it won't fail at runtime on a memory request but will run forever. But if multiple servers allocate too much memory at startup, then a new server trying to startup may fail. You need to configure the memory used on the command line. Go look at the advanced doc on command line parameters (non-server).

Commands affecting the server

Various `:xxxx` commands primarily control/affect the server.

:show serverlog

Toggles whether the server is logging data into the server logfile.

:show echoserver

If server logging is enabled, this will print the entries sent to the log on the print console as well.

:quit

This stops a running server, causing the program to exit. It will first flush any cached user files.

:restart

This will force the system to reload all its data files from disk (dictionary, topic data, live data) and then ask for your login. It's like starting the system from scratch, but it never stops execution. Good for revising a live server.

:crash

Force a crash to test system behavior during a crash.

:flush

If the server is caching user topic data in memory, this flushes all the cache items to disk.

Command Authorization

ChatScript has various `:xxxx` commands that can be given instead of normal chat input. These remain valid in server mode, and are a security liability if you have

potentially hostile users. Therefore all `:xxxx` commands require authorization (even in stand-alone mode). This is the file `authorizedIP.txt` existing in the top level directory. It will be read to validate any `:xxx` command. If the file is absent AND no `authorize=` command line parameter was given, authorization is granted.

The authorization file ships with a default `all`, WHICH IS DANGEROUS TO A SERVER. Normally the file will exist but have its content erased. That will lock out commands.

When you want to issue a command yourself to a running server, you go edit the file to insert on a line your own IP numeric address (since you have access to the file system itself, you are obviously an authorized person). Then you can log in to the ChatScript server and issue a `:xxxx` command.

The file can have any number of lines of IP addresses to be considered legal. You can also enter `L_login` lines, which will match a login name against the user's login. If that matches, that is also authorization. This works best when user names are assigned by a member authorization system, but it good enough in short durations if you use a weird enough login, e.g., `L_b1r8u2c9e0` (bruce interspersed with digits). Example file:

```
129.124.155.1
125.155.156.1
L_master
```

(a user named master case insensitive)

An authorization of `all` allows everyone. See also `authorize=` command line parameter and `nodebug` parameter.

CS as an embedded Client

Building a web-based interface to ChatScript is different from building an app involving ChatScript. In the web-based version, the ChatScript engine will always be a server on a reasonably powerful machine and have available the full resources of the web. In building an app-based NLP application with ChatScript, there are two general architectures, embed-based and server-based. Embedded based is described in the External communications manual.

Server-based

The other app configuration always requires an internet connection and communicates with a ChatScript server over it. In addition, this means it can easily do voice-to-text over it and text-to-voice back. All log files are already on the server so complete access to all conversations is conveniently available and all ChatScript data files are also easy to update as desired.

Of course the requirement of having an internet connection available is a severe one on the client and may also mean his app is generating data charges he may not like.

You still have a client app which sends and receives the communications (voices and/or texts) and must still decode return text from ChatScript to retrieve app commands for avatar behavior and other functions.

If you are using voice-to-text on the server, then you would create a server app that receives inputs from the client, communicates with the voice-to-text server, then hands appropriate text over to the ChatScript server.

The server app, on receiving the ChatScript output, may need to send it over to a text-to-voice conversion, and then ship all results back to the client. If you are not using voice in any capacity, the client app could communicate directly with the ChatScript server.

RESTful Server

ChatScript normally is not RESTful. It saves user state between volleys. But if you are scaling using multiple CS servers, you either need to arrange for routing specific users to specific CS servers where their files are, or have a central fileserver or database where user files are kept. And for proper full-scale service, you need a database that has multiple copies of data (like DynamoDB) that does not depend on a single machine. Or, your use of CS must be RESTful, meaning no user state is actually kept on the server side. This is partly possible. The data kept in the user topic file is:

1. prior things recently said by user and chatbot
2. user facts
3. state of every topic - which rules have been used up, changed topic flags
4. current topic & current interesting topic stack
5. input rejoinder tag
6. user variables & turn number and random seed
7. user context data - recently executed rule tags and on what volley

You can direct CS not to save ANY user data by turning on user caching, but setting the cache count to 0. Now you have a system that is RESTful, but can't keep any state. You can, however, send out state via out-of-band messages and read back in state the same way.

Currently there is a limit of 20k for an output message and user topic files vary in size including up to 40K (and you can allow them to be bigger via command line option). So it is not really practical to transmit all user state back and forth as oob data. Here is what is reasonable to accomplish:

1. Don't save most recent messages. You cannot detect the user repeating themselves automatically, though if you wrote out the most recent message you could manually check for immediate repeat.

2. Maybe write out user facts, depending on how many.
3. Don't change topic flags on the fly (so don't write them out). Don't write out the state of all rules (so you can't detect repeat or used up responders) but do write out the deepest gambit of a topic used. From that on start of a volley you can manually turn off all preceeding gambits.
4. Write out the current topic (for sure) and the interesting topic stack if you care.
5. Write out the rejoinder tag
6. Write out any user variables you care about. You may not care about the random seed or turn number, or you can write them out as well.
7. Currently there is no interface to context data so you can't write that out. Generally all this will be written out by a post-process phase.

Encryption

ChatScript normally reads and writes everything in plain text. If you need greater security then you need to do several things.

1. turn off server logs and user logs
2. Use `encrypt=` and `decrypt=` command line parameters.
3. Use default system or roll your own encryption in `privatecode`.

Encryption applies to user topic files and long term memory files (`^export` and `^import`). The built-in encryption method calls an external api server using `json` to perform encryption and decryption, passing the server urls provided by `encrypt=` and `decrypt=` and `ltmencrypt` and `userencrypt` command line parameters. The routines will add in the user's login id if you use `%s` in the command line parameters. Eg.,

```
encrypt=http://someapi/someotherdata/%s decrypt=http://someapi/%s/something
```

Where `%s` goes depends on the api you call.

To roll your own, you have to define routines in private code that `encrypt` and `decrypt` (see examples in `os.cpp`) and on private initialization override the default variables controlling `userFileSystem.encrypt` and `userFileSystem.decrypt`.

Client debug commands

When running a CS client talking to a cs server, there are some debugging commands that can prove useful. It is similar to local mode in that you have to log in with some arbitrary name. By default the bot goes first in the conversation, but if you want the human to go first, just put an `*` in front of your user name.

If you want to name the bot you want to talk to (in multibot servers), put a :botname after your login, e.g. `bruce:bot1`.

:source

This is the same :source command that you can use on a standalone server, sending the lines of the the file from client to server.

:dll filename

This loads the dll compile of CS and executes the commands of the file given, to test the dll.

:restart

Tell server to restart (reload) itself.

:converse

A conversation is a series of tab-delimited lines, where the first column is user name, the second is botname, and the third is user message. When user name changes, you have a new conversation.

JA internal commands

These commands read lines from a file (the first argument). If given, the second argument says how many lines to read before quitting. And the third argument, if given, tells how many lines to skip before really starting. The file is presumed to be at the top level of the ChatScript directory (if you dont give full path).

Typical data (except for :jmeter and :jaraw files) consists of tab or comma separated columns. First column is the user id. Second column is the category of the conversation. Third column is the specialty involved. Fourth column is the location, or some data to be passed back out unchanged. Sometimes this is converted status, or human labeled response or whatever. Fourth column is user input. Fifth column, if given, is the output expected from the bot, which is just echoed out.

45543099	Computer	Email	usa	Welcome! What's wrong with your email?
45543099	Computer	Email	usa	Could only receive
45543099	Computer	Email	usa	Are emails stuck in your outbox?
45543099	Computer	Email	usa	I have lost everything
45543100	Consumer_Electronics	TV	usa	Welcome! What's wrong with your TV?
45543100	Consumer_Electronics	TV	usa	White fuzzy
45543100	Consumer_Electronics	TV	usa	What's the brand and model number of your TV? H
45543100	Consumer_Electronics	TV	usa	2 years old

A full conversation starts with the bot remarks and alternates bot lines with user lines. Bot lines are there for completeness and are not passed to the server.

:jastarts filename {count {skip}}

Each line is an independent conversation of 1 line done by user. E.g.,

```
45543099    Computer    Email    usa Could only receive
45543100    Consumer_Electronics    TV    usa White fuzzy
```

Each conversation is started by sending a conversation initiation message of

```
[category: xxx specialty: yyy location: zzz id: iii expect: "..."]
```

where the location and expect fields only exists if the corresponding column was nonblank. Thereafter the lines sent to the server are just the user inputs.

:jaconverse file {count {skip}}

A conversation is a series of lines until a new user or category is detected. It starts with the opening message from the bot, and the lines alternate what the bot says (which are ignored and not passed to the server) and what the user says. E.g.,

:jamonologue file {count {skip}}

A conversation is a series of lines until a new user or category is detected. It consists ONLY of what the user says.

:jajmeter file {count {skip}}

In addition to file, if your file ends with /, that means use that as a directory and do all files within it (recursive).

A line is an entire conversation (a la JMETER std format), with what is expected from the bot and what the user says in return. An output file of TMP/jmeter.txt is generated which labels inputs with they passed (the output from bot matched expectation) or FAILED (and what oob data was shipped is shown along with what was expected and what was received).

Currently this only works for TCP, not api or integration which have different formats.

:jaraw file

Send the message under a [category: legal] header and the file only consists of input lines from user with no other data.